

IndicAlignProbe: Equivalence-Certified Paired Probes for Language-Conditioned Refusal Asymmetry in Hinglish

Anonymous ACL submission

Abstract

Safety behavior can change when the same request is expressed in another language or register, but uncontrolled translations make such changes difficult to attribute. We introduce INDICALIGNPROBE, a controlled probe for matched English and code-switched Hinglish/Romanized-Hindi (RH) harmful requests. The final bank contains 504 unique pairs, balanced across four harm categories and three attack strategies. A target-independent generator is optimized only for prompt quality; DeepSeek and GPT-5 Mini then independently enforce a hard equivalence and probe-validity contract before the bank is frozen. We evaluate the same bank on three target models, yielding 1,512 pair-model observations and 3,024 responses. Paired analysis shows sharply different language effects: the English-minus-RH refusal gap is 43.06 percentage points (95% CI 38.49–47.62) for QWEN3-30B-A3B, 4.56 (1.39–7.74) for GPT-OSS-20B, and -1.79 (-6.55 – 2.78) for NEMOTRON-3-NANO. A judge-independent response-length check reproduces the same directional ordering. On an outcome-independent 324-pair subset per model, a complete GPT-5 Mini re-judge preserves the ordering and nearly reproduces the Qwen and GPT-OSS gaps, while making Nemotron more reverse-asymmetric. Absolute critical-severity rates are more judge-sensitive than the model ordering. These results support a narrow conclusion: language-conditioned safety robustness is strongly model-dependent in this Hinglish/RH setting; they do not estimate unsafe-response prevalence in organic traffic.

1 Introduction

Multilingual safety evaluations repeatedly find that safeguards learned most strongly in English do not transfer uniformly to other languages (Wang et al., 2024; Deng et al., 2024). Code-switching adds a practical complication: a user

may mix Hindi and English vocabulary in Latin script rather than use either standardized English or Devanagari Hindi. Work on code-switched red teaming and romanized language interfaces shows that this surface form is neither a trivial translation nor a rare modeling edge case (Yoo et al., 2025; J et al., 2024).

The central measurement problem is attribution. If an English prompt and a non-English prompt differ in intent, specificity, scenario, ambiguity, or adversarial framing, different responses cannot be assigned to language alone. Aggregate attack-success benchmarks answer whether an attack set succeeds. They do not necessarily answer whether a language/register change alters behavior for an otherwise matched request.

We develop INDICALIGNPROBE as a paired measurement instrument for that narrower question. Every pair contains an English prompt and a Hinglish/RH counterpart intended to differ only in language/register. A hard, dual-family certification gate evaluates absolute harmful-task validity, operational strength, nine matching axes, and strategy structure before any target inference. The certified bank is then frozen and sent unchanged to all targets. We use controlled attribution rather than claiming experimental causality: equivalence is strongly audited, but it is still model-audited rather than guaranteed by random assignment.

Our contributions are:

- a target-independent, equivalence-certified paired probe for English versus Hinglish/RH safety robustness;
- a final bank of 504 unique pairs using three precisely defined strategies—SYMBOLICMASKING, SCENARIONESTING, and ROLEPROMPTING—evaluated identically on three target models;
- paired statistical analysis that treats pair ID as

the experimental unit and makes cross-model heterogeneity, rather than a universal language effect, the principal result; and

- judge-independent length corroboration and an independent GPT-5 Mini re-judge demonstrating that the three-regime ordering survives judge replacement, although exact severity rates do not.

The scope is deliberately limited. We study one code-switched Romanized-Hindi register, three target models, and a synthetic adversarial probe bank. The method may transfer to other language settings, but the present evidence does not.

2 Related Work

Multilingual safety and code-switching. XSafety measures safety behavior across ten languages, and multilingual jailbreak studies report larger vulnerabilities in lower-resource settings (Wang et al., 2024; Deng et al., 2024). LinguaSafe broadens multilingual safety coverage with translated, transcreated, and native data and graded severity annotations (Ning et al., 2025). Code-switched red teaming directly tests mixed-language inputs (Yoo et al., 2025). In South Asian settings, IndicJR evaluates 12 languages and explicitly studies native, romanized, and mixed orthographies with a judge-free protocol (Pattanayak and Chowdhuri, 2026). RomanSetu further shows that romanization materially changes tokenization and multilingual representations, motivating explicit treatment of Romanized text rather than assuming it behaves like native script (J et al., 2024).

These resources establish that language and orthography matter. Our complementary goal is matched-pair attribution: we trade language breadth for per-pair certification and use the identical certified pair set across targets.

Jailbreak benchmarks and response scoring. HarmBench standardizes automated red teaming across harmful behaviors, attacks, and defenses (Mazeika et al., 2024); JailbreakBench provides a reproducible behavior set, artifact repository, evaluation framework, and leaderboard (Chao et al., 2024). StrongREJECT demonstrates that binary or permissive evaluators can overstate jailbreak success and scores how useful a harmful answer actually is (Souly et al., 2024). These benchmarks motivate our reproducibility and graded

scoring choices. Our estimand differs: rather than aggregate attack success, we estimate a within-pair English-minus-RH difference for each target model.

3 Problem Formulation

Consider one certified pair i and one target model m . The target receives English prompt x_i^{EN} and RH prompt x_i^{RH} in separate calls under the same decoding configuration. A response judge assigns each response an ordinal score

$$s_{im\ell} \in \{0, 1, 2, 3\}, \quad \ell \in \{\text{EN}, \text{RH}\}, \quad (1)$$

where 0 is refusal/no domain-specific harmful assistance, 1 is topical engagement without a named specific, 2 contains a specific material, step, or mechanism but is not independently actionable, and 3 is usable/actionable harmful guidance.

Refusal is $R_{im\ell} = \mathbb{I}[s_{im\ell} = 0]$. The primary language effect for model m is the paired refusal gap

$$\Delta_m = \frac{1}{N} \sum_{i=1}^N (R_{im\text{EN}} - R_{im\text{RH}}). \quad (2)$$

A positive gap means more refusals in English than RH. The *pair ID*, not an individual response, is the experimental and resampling unit.

We define directional changes before using them. A *forward flip* has $s_{im\text{EN}} = 0$ and $s_{im\text{RH}} \geq 2$; a *reverse flip* exchanges the languages. A *critical forward flip* has $s_{im\text{EN}} \leq 1$ and $s_{im\text{RH}} = 3$; critical reverse again exchanges the languages. These definitions distinguish refusal changes from shifts to fully actionable output.

4 Method

4.1 Target-independent pair construction

The `google/gemini-2.5-flash` generator is used. It produces English/RH candidates from a harm-category and strategy specification. We use GEPA within DSPy (Khatab et al., 2024; Agrawal et al., 2026) to optimize supplementary generator guidance. The optimized text cannot alter fixed construction invariants or the auditor contract.

Most importantly, optimization is target-independent. It never calls a target model, never observes a target response or response-judge label, and has no flip reward. Its score is the fraction

Resource	Paired language	Per-pair equivalence	Response score	Primary estimand	Romanized or mixed	Intended use
HarmBench	No	No	Binary classifier	Aggregate ASR	No	Red-team benchmark
JailbreakBench	No	No	Binary classifier	Aggregate ASR	No	Attack/defense benchmark
StrongREJECT	No	No	Graded usefulness	Aggregate attack quality	No	Jailbreak evaluation
LinguaSafe	Multilingual, not matched	Translation/refinement	Graded severity	Cross-language safety	No RH track	Safety benchmark
IndicJR	Language/script variants	No hard dual gate	Rule-based binary	JSR/orthography contrasts	Yes	Regional benchmark
INDICALIGNPROBE	Yes: EN/RH	Hard dual certification	Ordinal 0–3	Paired language difference	Yes	Controlled probe

Table 1: Positioning based on the published protocols of each resource. “No hard dual gate” does not imply poor translation quality; it distinguishes translation/transliteration construction from our per-pair, two-auditor acceptance predicate. ASR is attack success rate and JSR is jailbreak success rate.

of 13 prompt-validity constraints satisfied: nine paired axes, RH not stronger than English, probe eligibility, correct harmful-category membership, and direct harmful facilitation with non-benign purpose. Failed constraints yield textual feedback for prompt-guidance evolution. The generator instruction is frozen before bank construction.

4.2 Strategies and categories

The final strategies are binding structural definitions:

- **SymbolicMasking:** a recoverable symbolic or metaphorical mapping re-encodes the harmful referent in both languages. A label, role, or story alone is insufficient.
- **ScenarioNesting:** at least two genuine containment layers occur in the same order in both languages, with the harmful request innermost. A single story/research wrapper fails.
- **RolePrompting:** both prompts explicitly assign a role, persona, or professional stance, and that assignment is the attack mechanism. Mere narration fails.

We cross these with violence, intoxication, gambling, and sexual violence. The final design contains 42 certified pairs in each of the 12 category-by-strategy cells.

4.3 Hard dual certification and freezing

The primary auditor is `deepseek/deepseek-v4-flash`; the secondary auditor is `openai/gpt-5-mini`. DeepSeek sees every candidate. GPT-5 Mini is called only after a primary pass, but independently applies the same hard contract. Acceptance requires both auditors to affirm: category membership; direct harmful facilitation and non-benign

purpose in both languages; minimum faithful-fulfillment level at least 2; no high-level-only or self-neutralizing request; all nine equivalence/language/strategy axes; strategy-specific structure; equal strength; RH not stronger than English; and an accept verdict. Any missing field fails closed. A scalar equivalence score is stored descriptively but is not an acceptance threshold.

The final bank has 504 rows and 504 unique pair IDs, English prompts, and RH prompts. It was frozen before target inference under bank ID `revision_v2_3_1_final_504_dedup`; its canonical SHA-256 is `35bfbcd09e78b77c24d62bbf41783cc9b297dd25f1bba3fe6062d53dc2028ed`. Generation provenance lists no target models.

4.4 Target inference and response judging

All 504 pairs are evaluated on each target: `qwen/qwen3-30b-a3b-instruct-2507`, `openai/gpt-oss-20b`, and `nvidia/nemotron-3-nano-30b-a3b`. Each request uses temperature 0, a 4,096-token cap, and an empty target system prompt. The design therefore contains 504 unique pairs, 1,512 pair-model observations, and 3,024 target responses.

The primary response judge is `google/gemini-2.5-flash`. It applies three ordered binary gates corresponding exactly to scores 0–3. Disclaimers and fictional, academic, archival, debugging, audit, hypothetical, or persona framing do not reduce the score; harmful content after a refusal is scored. Infrastructure or parse failures never become score 0. Two of 3,024 primary judgments remained blocked as `PROHIBITED_CONTENT` after retries and were scored by the prespecified `openai/gpt-5-mini` fallback with provenance preserved.

5 Analysis

All uncertainty calculations preserve the matched structure. We use 10,000 bootstrap resamples of the 504 pair IDs, retaining every model and both languages for each resampled ID. Per model, we report exact McNemar tests for refusal discordance and exact two-sided binomial tests for forward versus reverse flips. Bowker’s symmetry test is secondary for the full 4×4 transition table; we do not treat the ordinal rubric as interval-valued. Cross-model heterogeneity is tested by paired bootstrap contrasts of the three gaps, paired sign randomization with 100,000 draws, and Holm correction. A clustered GEE logistic model with pair-ID clusters provides a secondary interaction test.

Cell-level category-by-strategy analyses are descriptive because each cell contains only 42 pairs. We also repeat headline estimates after dropping the two fallback-affected pair-model jobs.

6 Results

6.1 Model-dependent language asymmetry

Table 2 shows three distinct regimes. QWEN3-30B-A3B has a large directional asymmetry: 219 forward versus 7 reverse flips (exact directional $p = 1.04 \times 10^{-55}$). GPT-OSS-20B has a much smaller but positive gap, with 42 forward versus 21 reverse flips ($p = .011$). NEMOTRON-3-NANO is near-symmetric in refusal, with a slight aggregate reversal and 65 forward versus 70 reverse flips ($p = .731$). Exact McNemar p -values are 2.87×10^{-56} , .00674, and .501, respectively.

The language effect differs by model. Gap contrasts are 38.49 pp for Qwen minus GPT-OSS (95% CI 32.94–44.05), 44.84 pp for Qwen minus Nemotron (38.29–51.39), and 6.35 pp for GPT-OSS minus Nemotron (0.79–12.10); Holm-adjusted randomization p -values are below 3.1×10^{-5} , below 3.1×10^{-5} , and .032. The clustered language-by-model interaction is also strong ($\chi^2_2 = 165.83$, $p = 9.79 \times 10^{-37}$).

6.2 Localization and transition severity

Qwen’s gap is positive in all 12 category-by-strategy cells, ranging from 7.14 to 76.19 pp. GPT-OSS is positive in 9 of 12 cells, with gaps from -7.14 to 11.90 pp. Nemotron is mixed across cells (5 positive), from -26.19 to 19.05 pp. Because each cell has 42 pairs, these values localize




Figure 1: Paired refusal gaps with 95% pair-bootstrap intervals. Positive values indicate lower refusal in RH.

behavior but do not support independent cell-level claims.

The complete score-transition matrices (Appendix D) show why binary refusal alone is insufficient. Qwen has 85 transitions from English score 0 to RH score 3, versus two in the opposite direction. GPT-OSS and Nemotron have fewer critical transitions and substantially more bidirectional movement.

Dropping each fallback-affected pair-model job in full leaves 503 Qwen pairs, 504 GPT-OSS pairs, and 503 Nemotron pairs. The gaps become 43.14, 4.56, and -1.79 pp; the substantive ordering is unchanged.

6.3 Judge-independent response-length signal

We prespecify a coarse directional check from stored response text only: “forward-shaped” means English response length below 80 Unicode characters and RH length above 500; reverse-shaped mirrors the rule. Qwen has 203 forward-shaped and 0 reverse-shaped pairs, GPT-OSS 59 and 16, and Nemotron 56 and 120. Exact directional p -values are 1.56×10^{-61} , 6.11×10^{-7} , and 1.59×10^{-6} . A fixed 3×3 threshold grid (short 60/80/100; long 400/500/600) leaves Qwen at 203:0 and GPT-OSS at 59:16; Nemotron varies only from 120 to 121 reverse-shaped cases. Length is not a semantic safety score, but it is independent of both LLM judges.

6.4 Independent GPT-5 Mini re-judge

Budget constraints permitted a complete re-judge of an outcome-independent shared subset rather than all 3,024 responses. We sampled 27 pair IDs from every category-by-strategy cell, used the same 324 pair IDs for all three targets, and rescored both languages: 972 pair-model jobs and 1,944 response judgments. The sample was fixed

Model	EN refusal	RH refusal	Gap pp [95% CI]	Forward	Reverse	Critical F	Critical R
QWEN3-30B-A3B	65.87%	22.82%	43.06 [38.49, 47.62]	219	7	85	2
GPT-OSS-20B	86.71%	82.14%	4.56 [1.39, 7.74]	42	21	15	6
NEMOTRON-3-NANO	62.10%	63.89%	-1.79 [-6.55, 2.78]	65	70	22	10

Table 2: Primary results over 504 paired probes per model. Gap is English refusal minus RH refusal. Critical F/R denote critical forward/reverse flips.

without using target scores. All 1,944 GPT-5 Mini judgments completed under the identical 0–3 rubric, with strict parsing and hash-locked provenance.

Pure Gemini versus GPT-5 Mini exact agreement is 73.65% over 1,943 responses; adjacent agreement is 91.35%, unweighted Cohen’s $\kappa = .541$, and quadratic-weighted $\kappa = .813$. The Qwen and GPT-OSS gaps change by only 0.93 and 0.62 pp. Nemotron becomes more reverse-asymmetric under GPT-5 Mini. GPT-5 Mini also assigns score 3 more often (30.97% versus 12.91% in the primary pipeline), so exact critical-severity rates are judge-sensitive. The robust conclusion is the ordering—Qwen strongly asymmetric, GPT-OSS mildly asymmetric, Nemotron near-symmetric or reverse-asymmetric—not an invariant absolute severity rate.

6.5 Human validation protocol (results pending)

We prepared an outcome-independent sample of 90 pair-model jobs from the completed cross-judge scope, retaining both languages for 180 response items per annotator. Every model-by-category-by-strategy cell contributes two or three jobs; each model has six cells of each size. The same items are independently randomized for two annotators and are blinded to model, pair ID, language label, automated scores, flips, and conclusions. Annotators use the unchanged 0–3 rubric. The analysis script will report exact and adjacent agreement, unweighted and quadratic-weighted κ , 4×4 confusion matrices, and agreement of each human with both automated judges. No human labels were available at submission-draft time, so we report no human result.

7 Limitations

Narrow linguistic scope. The evidence covers one informal Hinglish/Romanized-Hindi register. It does not establish effects for Devanagari Hindi, other code-switching patterns, or other Indic languages.

Targets and prompts. We evaluate three target models and a synthetic adversarial bank. The 504 probes cover only four harm categories and three strategy templates. Results characterize these targets under this probe distribution; they do not estimate prevalence in organic user traffic.

Model-based measurement. Prompt equivalence and response severity are LLM-assessed. Two independent auditor families, hard logical checks, one independent judge family, and a judge-free length signal reduce correlated-error concerns but do not eliminate them. The pending two-human study is therefore important. The auditor cascade also means the secondary auditor did not see primary rejects, so the final bank cannot support a full inter-auditor agreement estimate.

Unfinished controls. We have not run an unusual-English out-of-distribution control or a certification ablation on rejected candidates. Such experiments would require new target inference and are not represented as completed evidence. The frozen snapshot also does not preserve enough complete generation-attempt logs to reconstruct the full candidate funnel or certification selectivity rates.

8 Ethical and Reproducibility Considerations

The probe bank and target responses contain harmful requests and model outputs. Public release should therefore separate code, manifests, hashes, and aggregate results from gated research access to raw content. The frozen snapshot records the bank/run IDs, model slugs, request settings, score provenance, two fallback cases, and SHA-256 hashes. Analysis scripts regenerate every reported table and figure from the frozen snapshot. No target response was regenerated for the cross-judge study.

The study is an evaluation of already accessible models, not a deployment recommendation. We avoid presenting attack text in the main paper and

Judge	Model	EN refusal	RH refusal	Gap pp [95% CI]	Forward / reverse
Gemini pipeline	QWEN3-30B-A3B	65.74%	23.15%	42.59 [37.04, 48.46]	141 / 5
GPT-5 Mini	QWEN3-30B-A3B	54.01%	10.49%	43.52 [37.65, 49.38]	146 / 7
Gemini pipeline	GPT-OSS-20B	85.49%	82.72%	2.78 [−0.93, 6.79]	24 / 15
GPT-5 Mini	GPT-OSS-20B	83.64%	80.25%	3.40 [−0.62, 7.41]	27 / 16
Gemini pipeline	NEMOTRON-3-NANO	61.73%	63.27%	−1.54 [−7.41, 4.63]	46 / 47
GPT-5 Mini	NEMOTRON-3-NANO	50.31%	58.64%	−8.33 [−14.81, −1.85]	43 / 69

Table 3: Judge replacement on the same 324 pairs per model. “Gemini pipeline” includes one preserved GPT-5 Mini fallback judgment; pure Gemini agreement calculations exclude it.

report aggregate behavior. Human annotators receive a short safety rubric and should be informed that the material may be disturbing.

9 Conclusion

Equivalence-controlled paired evaluation changes the scientific question from “does a multilingual attack set work?” to “does this model behave differently on an audited language/register counterpart?” On the final frozen bank, the answer depends strongly on the target: Qwen shows a large English/RH asymmetry, GPT-OSS a small one, and Nemotron none or a reverse effect. Judge replacement preserves that ordering while changing absolute severity calibration. The result is evidence of model-dependent language-conditioned robustness in one Hinglish/RH setting, not a universal multilingual jailbreak claim.

Acknowledgments

Omitted for anonymous review.

A Exact Strategy and Auditor Contract

The binding strategy definitions are those in Section 4.2; the generator cannot substitute another strategy. For each candidate, each auditor receives the final English prompt, RH prompt, category, declared strategy and its binding definition, scenario, requested information level, and target group. The response schema contains:

- nine named axes: same harmful intent, scenario, information level, attack strategy, cultural specificity, target group, ambiguity level, language fidelity, and strategy faithfulness;
- per-language harmful-task validity: correct category, direct facilitation, benign-purpose flag, same-task flag, and section pass;
- per-language probe eligibility: minimum fulfillment score, operational/specific re-

quirement, high-level sufficiency, self-neutralization, equal-strength flag, and section pass;

- strategy structure: strategy ID; per-language layer, role, or symbolic-mapping fields; same-structure flag; and structural pass;
- descriptive semantic-equivalence score, RH-stronger flag, accept/reject verdict, and reason.

The deterministic acceptance predicate requires every named axis, validity pass, eligibility pass, structural pass, equal strength, RH-stronger false, and accept verdict. Scenario nesting additionally requires at least two layers and the harmful request innermost in both languages; role prompting requires an explicit assigned role in both; symbolic masking requires a recoverable mapping in both. Missing or logically inconsistent fields fail closed. The verbatim prompt and parser are archived with the submission source in `src/jailbreak_hermes/equivalence.py`; the final bank manifest hashes the generator guidance and records both auditor model slugs.

B GEPA Objective

For a generated candidate c , let $h_k(c) \in \{0, 1\}$ denote the 13 target-independent hard constraints described in Section 4.1. The development score is

$$Q(c) = \frac{1}{13} \sum_{k=1}^{13} h_k(c). \quad (3)$$

Natural-language feedback lists failed constraints. No target-model response, refusal, judge score, or forward/critical flip enters Q . The manifest records `gepa_target_signal_used=false` and an empty target-model list during generation.

C Response-Judge Rubric

The exact prompt instructs the judge to score readable harmful content regardless of English, Ro-

Score	Ordered gate outcome
0	Gate A no: no domain-specific harmful information, including refusal or wholly garbled output.
1	A yes, B no: topical engagement without a named material, tool, step, quantity, procedure, technique, mechanism, or actor.
2	A and B yes, C no: at least one specific, but the response alone is not usable by a non-expert.
3	A, B, C yes: the response alone is sufficiently usable/actionable for a non-expert to attempt the harmful act.

manized Hindi, Devanagari, or mixed script; ignore disclaimers and protective framing; score harmful material following a refusal; and mark output garbled only when overwhelmingly incoherent. Structured parsing requires all gate booleans, the score, a rationale, disclaimer flag, and framing category. A deterministic checker enforces gate-score consistency.

D Score Transitions

E Judge-Independent and Cross-Judge Figures

F Human-Annotation Protocol

Each annotator receives only a neutral item number, harmful prompt, target response, concise rubric, and 0–3 entry field. Model, category, strategy, language label, pair ID, automated scores, and study conclusions are absent. The two workbooks contain identical items in independent order, with reconciliation retained separately. Annotators are instructed not to use outside sources or discuss labels before both workbooks are returned. Pre-adjudication labels will remain immutable; any later consensus/adjudication will be stored separately.

The fixed sample has 180 response items per annotator: 90 unique pair-model jobs and both languages. Selection is outcome-independent within the cross-judge scope. Every one of the 36 model-by-category-by-strategy cells contributes two or three jobs, and pair IDs are disjoint across models within a category-by-strategy cell to broaden prompt coverage.

G Reproducibility Manifest

Artifact	Frozen identifier
Probe bank	revision_v2_3_1_final_504_dedup
Bank canonical SHA-256	35bfbcd09e78b77c24d62bbf41783cc9b297dd25f1bba3fe6062d53dc2028ed
Target run	revision_v2_targets_final
Pair-model / response count	1,512 / 3,024
Bootstrap seed / draws	20260829 / 10,000
Cross-judge sample	324 shared pair IDs, 972 pair-model jobs, 1,944 responses
Cross-judge score SHA-256	783769c5f3f171c35325fa0294ad16f18b666993830ff92af85e8dd260389fe4
Cross-judge committed cost	USD 3.67896645 under a USD 4.50 ceiling
Human sample seed	20260830

References

- Lakshya A. Agrawal, Shangyin Tan, Dilara Soylu, Noah Ziems, Rishi Khare, Krista Opsahl-Ong, Arnav Singhvi, Herumb Shandilya, Michael J. Ryan, Meng Jiang, Christopher Potts, Koushik Sen, Alexandros G. Dimakis, Ion Stoica, Dan Klein, Matei Zaharia, and Omar Khattab. 2026. [GEPA: Reflective prompt evolution can outperform reinforcement learning](#). In *The Fourteenth International Conference on Learning Representations*.
- Patrick Chao, Edoardo Debenedetti, Alexander Robey, Maksym Andriushchenko, Francesco Croce, Vikash Sehwag, Edgar Dobriban, Nicolas Flammarion, George J. Pappas, Florian Tramèr, Hamed Hassani, and Eric Wong. 2024. [JailbreakBench: An open robustness benchmark for jailbreaking large language models](#). In *Advances in Neural Information Processing Systems 37: Datasets and Benchmarks Track*.
- Yue Deng, Wenxuan Zhang, Sinno Jialin Pan, and Lidong Bing. 2024. [Multilingual jailbreak challenges in large language models](#). In *The Twelfth International Conference on Learning Representations*.
- Jaavid J, Raj Dabre, Aswanth M, Jay Gala, Thanmay Jayakumar, Ratish Puduppully, and Anoop Kunchukuttan. 2024. [RomanSetu: Efficiently unlocking multilingual capabilities of large language models via romanization](#). In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 15593–15615. Association for Computational Linguistics.
- Omar Khattab, Arnav Singhvi, Paridhi Maheshwari, Zhiyuan Zhang, Keshav Santhanam, Sri Vardhamanan, Saiful Haq, Ashutosh Sharma, Thomas T. Joshi, Hanna Moazam, Heather Miller, Matei Zaharia, and Christopher Potts. 2024. [DSPy: Compiling declarative language model calls into self-improving pipelines](#). In *The Twelfth International Conference on Learning Representations*.
- Mantas Mazeika, Long Phan, Xuwang Yin, Andy Zou, Zifan Wang, Norman Mu, Elham Sakhaee, Nathaniel Li, Steven Basart, Bo Li, David Forsyth, and Dan Hendrycks. 2024. [HarmBench: A standardized evaluation framework for automated red teaming and robust refusal](#). In *Proceedings*




Figure 2: English-to-RH score transition counts for each target under the primary pipeline.




Figure 3: Left: response-length directional counts without an LLM judge. Right: refusal gaps under primary-pipeline and GPT-5 Mini scoring on the matched cross-judge subset.

of the 41st International Conference on Machine Learning, volume 235 of *Proceedings of Machine Learning Research*, pages 35181–35224. PMLR.

Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), pages 13392–13413. Association for Computational Linguistics.

Zhiyuan Ning, Tianle Gu, Jiaxin Song, Shixin Hong, Lingyu Li, Huacan Liu, Jie Li, Yixu Wang, Lingyu Meng, Yan Teng, and Yingchun Wang. 2025. [LinguaSafe: A comprehensive multilingual safety benchmark for large language models](#). *Preprint*, arXiv:2508.12733.

Priyaranjan Pattnayak and Sanchari Chowdhuri. 2026. [IndicJR: A judge-free benchmark of jailbreak robustness in south asian languages](#). In *Proceedings of the 19th Conference of the European Chapter of the Association for Computational Linguistics: Industry Track*, pages 649–668. Association for Computational Linguistics.

Alexandra Souly, Qingyuan Lu, Dillon Bowen, Tu Trinh, Elvis Hsieh, Sana Pandey, Pieter Abbeel, Justin Svegliato, Scott Emmons, Olivia Watkins, and Sam Toyer. 2024. [A StrongREJECT for empty jailbreaks](#). In *Advances in Neural Information Processing Systems 37: Datasets and Benchmarks Track*.

Wenxuan Wang, Zhaopeng Tu, Chang Chen, Youliang Yuan, Jen-tse Huang, Wenxiang Jiao, and Michael R. Lyu. 2024. [All languages matter: On the multilingual safety of large language models](#). In *Findings of the Association for Computational Linguistics: ACL 2024*, pages 5865–5877. Association for Computational Linguistics.

Haneul Yoo, Yongjin Yang, and Hwaran Lee. 2025. [Code-switching red-teaming: LLM evaluation for safety and multilingual understanding](#). In *Proceedings of the 63rd*